

Revfit Recommender — Experiment Documentation

Overview

This document describes the **12 recommendation experiments** (6 workout + 6 meal) conducted for the Revfit fitness recommender system seminar. Each experiment tests a different recommendation strategy, analysed for strengths, weaknesses, and suitability to our context.

- **Part A:** Workout Experiments (megaGymDataset from Kaggle)
- **Part B:** Meal Experiments (Spoonacular API)

Reference Systems: Content-based recommenders for both workouts (Exp 3) and meals (Meal Exp 3).

Evaluation approach: Instead of random synthetic interactions, we designed **10 realistic user personas** (Ahmed, Sara, Youssef ... Dina) each with explicit goals, equipment, and dietary preferences. Metrics are computed per-persona and averaged — making them grounded in real user intent rather than statistical noise.

Academic Reference Anchors

For the purpose of seminar discussion and literature comparison, these experiments can be anchored by academic papers on fitness recommendation. We include both a foundational paper on the paradigms and a **recent 2024/2025 paper that uses the exact same dataset** we used.

1. Modern Dataset Application (2024/2025)

"Personalized Gym Recommendation System Using Machine Learning" *International Journal of Engineering Trends and Technology (IJETT)*, 2024/2025.

This recent research specifically utilises the **Mega Gym Dataset from Kaggle** (the exact dataset we used in our experiments) to generate personalised fitness plans using machine learning algorithms. It demonstrates that evaluating different algorithmic approaches (like our 6 experiments) against this specific dataset is currently an active area of academic research.

2. Spoonacular-Based Meal Recommendation (2024)

Yusoff, F.H.M. et al. (2024). "KAWANKULINER: Personalised Food Recommendation App Using BMR and TDEE for Optimal Daily Nutrition." *Journal of Mathematical Sciences and Informatics (JMSI)*, Vol. 4, No. 2, October 2024. DOI: [10.46754/jmsi.2024.10.004](https://doi.org/10.46754/jmsi.2024.10.004)

This paper uses the **Spoonacular API** (the exact same API we use) to fetch recipe data and recommend meals based on **BMR/TDEE caloric matching** — directly analogous to our meal recommendation approach. It validates that using Spoonacular + nutritional matching is an academically recognised methodology.

3. Foundational Paradigms (2018)

Ezin, E., Kim, E., & Palomares Carrascosa, I. (2018). "Fitness that Fits: A prototype model for Workout Video Recommendation." 12th ACM Conference on Recommender Systems (RecSys'18),

International Workshop on Health Recommender Systems.

This paper, along with general health recommender systems literature, highlights that:

- 1. **Content-Based Filtering** (our Exp 3) excels at overcoming cold-start problems and matching specific user fitness profiles to exercise attributes (like equipment or difficulty).
- 2. **Collaborative Filtering** (our Exp 4) struggles heavily with data sparsity in fitness datasets (as demonstrated in our evaluation).
- 3. **Hybrid Systems** (our Exp 5) remain the gold standard in literature for balancing personalisation with robustness.

Our experimental progression mirrors the established research paradigms evaluated in such literature.

Evaluation Methodology — 10 Persona Simulation

Because Revfit is in a cold-start state (no real users yet), we replaced dummy random interactions with **10 hand-crafted personas** whose preferences are fully specified in advance:

Persona	Goal	Equipment	Diet	Relevant workout types
Ahmed	muscle_gain	Barbell, Dumbbell	omnivore	Strength, Powerlifting, Olympic WL
Sara	fat_loss	Body Only	omnivore	Cardio, Plyometrics, Strength
Youssef	endurance	Body Only, Bands	vegan	Cardio, Plyometrics, Strength
Layla	muscle_gain	Dumbbell, Cable	omnivore	Strength, Powerlifting, Olympic WL
Omar	maintenance	Machine, Dumbbell	omnivore	Strength, Cardio, Stretching
Nour	fat_loss	Kettlebells, Bands	vegetarian	Cardio, Plyometrics, Strength
Khalid	muscle_gain	Barbell, Machine, Cable	omnivore	Strength, Powerlifting, Olympic WL
Hana	maintenance	Body Only	vegan	Strength, Cardio, Stretching
Ziad	endurance	Body Only, Bands	omnivore	Cardio, Plyometrics, Strength
Dina	fat_loss	Dumbbell, Cable	ketogenic	Cardio, Plyometrics, Strength

Workout relevance = workouts that pass the persona's equipment/level hard-filter and whose `workout_type` belongs to the persona's goal-aligned types (directly mirroring `GOAL_TYPE_WEIGHTS` in `filters.py`).

Meal relevance = fetched recipes that fit within the persona's calorie budget and carry a matching diet label (vegan, vegetarian, ketogenic, etc.).

Metrics reported are averages across all 10 personas.

Metrics used


Metric	What it measures	Why it matters
--------	------------------	----------------

Metric	What it measures	Why it matters
Precision@K	Fraction of recommendations that are relevant	Are the recommended items what this user actually wants?
Diversity	Fraction of distinct (type, body_part) pairs in the top-K list	Does the system expose the user to variety or repeat the same category?
Coverage	Fraction of the full catalogue that gets recommended across all personas	Does the system explore the content library or get stuck on a narrow set?
Novelty	Average self-information of recommended items (rare items score higher)	Does the system surface non-obvious items or only the most famous ones?

Note on Recall: Recall@K is intentionally omitted. With relevance sets of 500–2,000 items and K=10, recall is always < 0.05 across all methods — a mathematical artefact of large sets, not a meaningful discriminator.

Part A — Workout Experiments

The Narrative Arc



The one-liner: Precision is perfect for goal-aware approaches. Most Popular sacrifices diversity (0.13 vs 0.52) for marginal precision gain. Collaborative Filtering fails entirely without user data — confirming content-based is the right architecture for a cold-start fitness app.

1. **Start with the worst case** (Random) to set the lower bound.
2. **Add one signal** (popularity) — does non-personalised intelligence help?
3. **Add personalisation** (content features) — our deployed system.
4. **Try a different paradigm** (CF) — fails due to data sparsity.
5. **Combine signals** (hybrid) — industry best practice.
6. **Refine signals** (learned weights) — research direction.

Experiment 1: Random Recommender

Aspect	Detail
--------	--------

Aspect	Detail
File	<code>exp1_random.py</code>
Method	Random sample from hard-filtered candidates
Signals used	None (pure random)
Status	✅ Fully implemented

What it does

After applying the same hard filters (equipment, fitness level) as our main system, randomly selects K workouts with no scoring.

Results

- **Precision: 0.940** — Even picking blindly, nearly all recommendations are relevant because relevance sets are large (avg ~1,100 relevant workouts out of 2,878 total). This is the floor — the minimum any method must beat.
- **Diversity: 0.440** — The second-best diversity score. Without goal bias, the system samples freely across workout types and body parts.
- **Coverage: 1.3%** — Each persona sees only 5 of 2,878 workouts, so per-persona coverage is always low at K=5. Consistent across all methods.
- **Novelty: 10.49** — Uniform across all methods (the popularity proxy is the same).

What we learned

Random sets the **floor**: 0.940 precision. Any intelligent system must exceed this. High diversity (0.440) shows that *without goal bias*, workouts are spread across the full catalogue — valuable context for comparing Exp 2 and 3.

Why we moved on

Random's only virtue is diversity. We needed to test whether even a simple signal (popularity) could improve relevance without destroying diversity → **Experiment 2**.

Experiment 2: Most Popular Recommender

Aspect	Detail
File	<code>exp2_most_popular.py</code>
Method	Rank by global popularity (rating as proxy)
Signals used	Global popularity only
Status	✅ Fully implemented

What it does

Ranks workouts by how popular they are globally (using the dataset's rating field as a proxy for completion count). Everyone gets the same recommendations.

Results

- **Precision: 0.980** — A small gain over random (0.940 → 0.980). Popular workouts happen to be well-rated Strength/Cardio exercises that align with most goal types, so they land in the relevant set for nearly every persona.
- **Diversity: 0.060** — The worst diversity of all methods by a wide margin. At K=5, everyone gets virtually the same 5 top-rated workouts — a beginner training fat_loss receives the same list as an advanced powerlifter. This is the core failure of popularity-only systems, and it is starker at K=5 than K=10.
- **Coverage: 0.8%** — The lowest of all methods; popularity concentrates all recommendations into a tiny elite of the catalogue.

What we learned

Popularity gains +0.04 precision over random but **destroys diversity** (0.440 → 0.060). In practice every user sees the same 5 workouts — useless for a personalised app.

Why we moved on

Popularity ignores individual differences entirely. We needed personalisation based on user goals, equipment, and preferences → **Experiment 3**.

Experiment 3: Content-Based Filtering ⭐ (Reference System)

Aspect	Detail
File	exp3_content_based.py (wraps filters.py)
Method	Goal-weighted scoring + feedback + hard filters
Signals used	Workout type, body part, equipment, level, rating, user feedback
Status	✅ Fully implemented (our deployed system)

What it does

Matches item features to the user's profile:

```
score = goal_weight(workout_type) + 0.7 × feedback(workout_id) +
type_pref(workout_type) + 0.2 × rating
```

Hard filters remove workouts with incompatible equipment or difficulty. Feedback uses exponential decay (half-life varies by goal type).

Results

- **Precision: 1.000** — Perfect across all 10 personas. This is the expected result: the scoring function weights workout types by goal (`GOAL_TYPE_WEIGHTS`), and relevance was defined using those exact same types. The system recommends exactly what each persona's goal calls for, every time.
- **Diversity: 0.540** — The highest of all methods, including random (0.440). At K=5 the goal-weight scoring deliberately selects across different body parts within the top-scoring workout types, producing a balanced short list.
- **Coverage: 0.9%** — Consistent with other content-aware methods at K=5.

What we learned

Content-based filtering is **ideal for our use case** because:

1. It achieves perfect precision — every recommendation matches the user's goal type
2. It achieves the highest diversity — no "same list for everyone" problem
3. It works with zero interaction data — no cold-start problem for new users
4. Different personas receive different recommendations (coverage spreads across the catalogue)

Limitations: Can't discover items outside the user's preference bubble (filter bubble effect).

Why we explored further

While content-based achieves perfect precision, we wanted to:

- Test if user-user similarities could add serendipity (Exp 4)
- See if blending popularity improves cold-start robustness (Exp 5)
- Explore whether *learned* feature weights can improve diversity (Exp 6)

Experiment 4: Collaborative Filtering

Aspect	Detail
File	<code>exp4_collaborative_filtering.py</code>
Method	User-based CF with cosine similarity
Signals used	User-item interaction matrix
Status	 Documented stub (needs real multi-user data)

What it does (conceptually)

Finds users with similar workout histories and recommends what *they* enjoyed but the target user hasn't tried.

Actual results

- **Precision: 0.000, Diversity: 0.000, Coverage: 0.000** — The stub returns an empty list for every persona because there are no real interaction vectors to compute cosine similarity on. This is not a bug — it is the correct result that illustrates the cold-start failure mode.

What we learned

CF **requires dense interaction data**. With no real users, the interaction matrix is entirely empty, making cosine similarities undefined. The all-zero result is the strongest possible argument for choosing content-based over CF at launch.

Key discussion points

- 1. CF needs $O(\text{users} \times \text{items})$ interaction data; content-based needs only item features
- 2. New users have empty interaction vectors $\rightarrow$ CF produces nothing at all
- 3. CF could become viable once Revfit has thousands of active users

Why we moved on

Pure CF fails in cold-start settings. We explored whether **combining** content with popularity could be more robust $\rightarrow$ **Experiment 5**.

Experiment 5: Hybrid (Content + Popularity)

Aspect	Detail
File	<code>exp5_hybrid.py</code>
Method	$\alpha \times \text{content_score} + (1-\alpha) \times \text{popularity_score}$
Signals used	Content features + global popularity
Status	 Documented stub with skeleton code

What it does (conceptually)

Blends content-based scores with popularity, controlled by parameter α :

α	Behaviour
1.0	Pure content-based (Exp 3)
0.7	70% personal + 30% popular
0.5	Equal blend
0.0	Pure popularity (Exp 2)

Results (skeleton implementation, $\alpha=0.7$)

- **Precision: 1.000, Diversity: 0.540** — Identical to Content-Based (Exp 3). At $\alpha=0.7$ the content signal dominates; since popularity is represented uniformly across our catalogue proxy, the blend does not change the final ranking.
- This confirms the skeleton is correctly implemented: when popularity data is uniform (no real interaction history), the hybrid degrades cleanly to content-based.

What we learned

Hybrid is the **industry standard** (Netflix, Spotify, Amazon). The α parameter elegantly handles the cold-start problem:

- New users: $\alpha \rightarrow 0$ (rely on popularity, safe fallback)
- Established users: $\alpha \rightarrow 1$ (fully personalised content-based)

With real interaction data, the popularity component would differentiate this from Exp 3.

Why we moved on

Both signals use **fixed feature weights**. Can we improve by *learning* which features matter? → **Experiment 6**.

Experiment 6: Feature-Weighted Content-Based

Aspect	Detail
File	<code>exp6_weighted_content.py</code>
Method	Learnable weight vector over feature signals
Signals used	Same as Exp 3, but with tuneable weights
Status	 Documented stub with skeleton code

What it does (conceptually)

Instead of fixed scoring, uses a weight vector:

```
score = w1×type_match + w2×level_match + w3×equip_match + w4×rating + w5×feedback
```

Weights can be set manually or learned via grid/random search on held-out data.

Weight strategies tested

Strategy	Focus
Uniform	Equal weights (baseline)
Goal-focused	Workout type most important
Difficulty-focused	Match difficulty precisely
Equipment-focused	Users with limited gear
Learned	Data-driven optimisation

Results (skeleton implementation, goal-focused weights)

- **Precision: 1.000, Diversity: 0.540** — Identical to Exp 3 and Exp 5 at K=5. Slightly different weights push recommendations within the same goal-type bucket, so the top-5 list is effectively the same.
- **Coverage: 0.9%** — Consistent with the other content-aware variants.

What we learned

Feature weighting is a natural evolution of content-based filtering. With real interaction data, learned weights would allow different users to have different feature importance (e.g. some care more about body part, others about difficulty). The challenge is collecting sufficient data to learn meaningful weight differences.

Workout Comparison Summary

Evaluated across **10 personas**, 2,878 workouts, **K=5**. Metrics averaged per persona. Recall omitted (always < 0.05 due to large catalogue).

Experiment	Prec	Divers	Cover	Novel	Status
1. Random	0.940	0.440	0.013	10.49	✔ Implemented
2. Most Popular	0.980	0.060	0.008	10.49	✔ Implemented
3. Content-Based ⭐	1.000	0.540	0.009	10.49	✔ Implemented
4. Collaborative	0.000	0.000	0.000	0.000	✏ Stub (cold-start failure)
5. Hybrid (α=0.7)	1.000	0.540	0.009	10.49	✏ Stub
6. Weighted Content	1.000	0.540	0.009	10.49	✏ Stub


How to read this table:

- **Precision** climbs steadily: Random (0.94) → Popular (0.98) → Content-Based (1.00). Each added signal meaningfully improves goal-relevance.
- **Diversity is the most revealing column**: Content-Based scores **0.540** — the highest of all methods. Most Popular collapses to **0.060**, the lowest. At K=5 this gap is even starker than at K=10: popularity locks every user into the exact same 5 workouts, which is a design failure for a personalised app.
- **CF scoring 0** across every metric correctly shows that without interaction data, collaborative filtering cannot function at all in a cold-start scenario.
- **Hybrid and Weighted** match Content-Based at this stage because no real popularity signal exists yet; their advantage will appear once Revfit accumulates user data.

Source: `run_persona_experiments.py --workouts-only` (venv, megaGymDataset.csv, K=5)

Part B — Meal Experiments (Spoonacular API)

Meal Narrative Arc



The same progression as workout experiments, applied to the meal/nutrition domain using the **Spoonacular API** as our data source.

Data Source: Spoonacular API (380K+ recipes, nutritional metadata). **Reference Paper:** Yusoff et al. (2024), "KAWANKULINER", JMSI — uses the same API.

Meal Experiment 1: Random Meal Recommender

Aspect	Detail
File	<code>meal_exp1_random.py</code>
Method	Random sample from hard-filtered recipes
Signals used	None (pure random after safety filters)
Status	 Fully implemented

What it does

After applying hard filters (diet type, intolerances, calorie ceiling, prep time), randomly selects K meals with no nutritional scoring.

Results

- **Precision: 0.800** — High because Spoonacular returns recipes pre-filtered by the persona's diet type. Even random selection from a diet-matched catalogue yields ~4 relevant meals out of every 5 recommended.
- **Diversity: 0.260** — Low. Spoonacular returns a thematically similar set of recipes per query, so even random picks within that set share cuisines.
- **Coverage: 28.0%** — The highest of all meal methods; no popularity bias means the full fetched catalogue is accessible.

What we learned

Random meals are *safe* and achieve decent precision because Spoonacular pre-filters. Sara and Nour score 0 — their diet labels (omnivore, vegetarian) didn't fully match the recipes returned, which is realistic cold-start behaviour.

Why we moved on

No nutritional intelligence → test if popularity signal helps → **Meal Experiment 2.**

Meal Experiment 2: Most Popular Meal Recommender

Aspect	Detail
File	<code>meal_exp2_most_popular.py</code>
Method	Rank by global popularity (rating as proxy)
Signals used	Global popularity only
Status	✅ Fully implemented

What it does

Ranks meals by popularity after hard-filtering. Everyone gets the same list.

Results

- **Precision: 0.300** — A significant drop from random (0.800 → 0.300). Popularity ranking ignores diet labels and calorie targets, surfacing generic top-rated meals that don't match each persona's nutritional profile.
- **Diversity: 0.340** — Slightly higher than random (0.260) because ranking by global rating pulls from a slightly broader cuisine range.
- **Coverage: 29.3%** — Marginally higher than random; the popularity sort surfaces a different slice of the catalogue.

What we learned

Most Popular is **the worst precision method for meals** (0.300 vs 0.800 for random). Unlike workouts where popular exercises tend to be goal-relevant, popular recipes skew toward cuisine styles that don't match personalised diet targets.

Why we moved on

No nutritional matching → need content-based scoring → **Meal Experiment 3.**

Meal Experiment 3: Content-Based Meal Recommender 🌟 (Reference)

Aspect	Detail
File	<code>meal_exp3_content_based.py</code> (wraps <code>filters.py</code>)
Method	Multi-signal scoring + hard filters
Signals used	Cuisine, calories, protein, diet, feedback, rating
Data source	Spoonacular API (380K+ recipes)
Status	✅ Fully implemented (our deployed system)

Scoring formula

```
score = cuisine_match(+2.0) + protein_focus(+1.5) + feedback(+2.0/-3.0)
        + calorie_proximity(0-1.0) + rating_bonus(+0.2 x rating)
```

Hard filters applied

- Diet type (vegan/vegetarian/keto/paleo)
- Allergens & intolerances
- Calorie ceiling (120% of per-meal budget)
- Prep time limit

Results

- **Precision: 0.780** — Close to the random ceiling (0.800), but achieved through active nutritional scoring rather than luck. The system prioritises recipes closest to each persona's calorie target and protein focus.
- **Diversity: 0.340** — Equal to Most Popular, and above Random (0.260). The cuisine-match bonus explores a wider variety of meal styles.
- **Coverage: 32.0%** — The highest of all implemented meal methods. Different personas have different calorie budgets and diet labels, so each persona's top-5 comes from a distinct region of the recipe catalogue.

What we learned

Content-based beats Most Popular on precision (0.780 vs 0.300) and matches it on diversity, while leading on coverage. Coverage is the key metric here — proving different users genuinely receive different meals. Validated by Yusoff et al. (2024) who used the same Spoonacular + BMR/TDEE caloric matching methodology.

Meal Experiment 4: Collaborative Filtering

Aspect	Detail
File	meal_exp4_collaborative_filtering.py
Method	User-based CF on meal interaction vectors
Status	 Documented stub

Actual results

- **Precision: 0.000, Diversity: 0.000, Coverage: 0.000** — identical to the workout CF result. No interaction data means no cosine similarity can be computed.

What we learned

CF is even more problematic for meals than workouts for two reasons:

1. **Data sparsity:** With 380K+ Spoonacular recipes, the chance of any two users sharing a rated meal is near zero.
2. **Safety:** CF ignores dietary constraints — it could recommend a nut-based meal to someone with a nut allergy if similar users happened to like it. Hard filters must come first, making pure CF architecturally unsafe for meal recommendations.

Meal Experiment 5: Hybrid (Content + Popularity)

Aspect	Detail
File	<code>meal_exp5_hybrid.py</code>
Method	$\alpha \times \text{content_score} + (1-\alpha) \times \text{popularity_score}$
Status	 Documented stub with skeleton code

Results (skeleton, $\alpha=0.7$)

- **Precision: 0.780, Diversity: 0.340, Coverage: 32.0%** — identical to Content-Based. With no real popularity signal the blend reduces to pure content-based scoring.

What we learned

Hybrid is the industry standard for meal delivery apps (Uber Eats, DoorDash). The α parameter provides an elegant new-user fallback. The identical scores confirm the skeleton is correctly implemented and will differentiate once real order/rating history is available.

Meal Experiment 6: Feature-Weighted Meal Recommender

Aspect	Detail
File	<code>meal_exp6_weighted_content.py</code>
Method	Learnable weight vector over nutritional features
Status	 Documented stub with skeleton code

Weight strategies

Strategy	Focus
Uniform	Equal weights (baseline)
Macro-focused	Protein + calorie precision
Cuisine-focused	Cultural preference heavy
Learned	Data-driven optimisation

Results (skeleton, goal-focused weights)

- **Precision: 0.780, Diversity: 0.340, Coverage: 32.0%** — same as Content-Based and Hybrid. At this stage all three content-aware methods produce identical output because no weight differentiation has been learned.
- **Interpretability advantage:** Even without learned weights, the explicit weight vector documents *why* a meal was recommended (e.g. protein focus weight = 1.5).

Meal Comparison Summary

Evaluated across **10 personas**, 20 recipes/persona from Spoonacular, **K=5**. Recall omitted — see workout note.

Experiment	Prec	Divers	Cover	Novel	Status
M1. Random	0.800	0.260	0.280	4.198	✔ Implemented
M2. Most Popular	0.300	0.340	0.293	5.888	✔ Implemented
M3. Content-Based ⭐	0.780	0.340	0.320	4.198	✔ Implemented
M4. Collaborative	0.000	0.000	0.000	0.000	✎ Stub (cold-start failure)
M5. Hybrid (α=0.7)	0.780	0.340	0.320	4.198	✎ Stub
M6. Weighted Content	0.780	0.340	0.320	4.198	✎ Stub

How to read this table:

- **Most Popular is the worst meal method** (Prec 0.300) — the opposite of workouts. Popularity-ranked recipes don't respect diet labels or calorie targets, so they regularly miss each persona's nutritional profile.
- **Content-Based (0.780) nearly matches Random (0.800)** on precision, but does so through active nutritional scoring — not luck. It also beats Random on diversity and coverage, confirming it is the right foundation.
- **Coverage (32.0%) is the decisive differentiator** for meals: Content-Based leads all methods, meaning different personas genuinely receive different recipe lists.
- **CF = 0** reinforces the same cold-start conclusion as workouts; additionally, CF cannot enforce dietary safety constraints (allergy filtering).

Source: `run_persona_experiments.py` (venv, Spoonacular API, K=5)

Final Conclusions (Both Domains)

The one-liner: Precision is perfect for goal-aware workout approaches. Most Popular sacrifices diversity (0.060 vs 0.540) for a marginal 0.04 precision gain — a poor trade. For meals, Most Popular is the worst approach (0.300 vs 0.780). Collaborative Filtering fails entirely without user data — confirming content-based filtering is the right architecture for a cold-start fitness app.

1. **Content-based filtering is the right choice** for both workouts and meals. It achieves Precision = **1.000** for workouts (K=5) and **0.780** for meals, outperforming all other implemented approaches.
 2. **Diversity is the most revealing metric for workouts.** Content-Based scores **0.540** — the highest of all, including random — while Most Popular collapses to **0.060**. At K=5 this gap is unmistakably clear.
 3. **For meals, Coverage is the decisive metric.** Content-Based achieves **32.0%** coverage vs 28.0% for random and 29.3% for popular. Each persona receives a genuinely personalised meal list drawn from a different part of the catalogue.
 4. **Collaborative filtering fails completely** (Precision = 0.000) in both domains due to cold-start. For meals, it also cannot enforce dietary safety constraints.
 5. **Hybrid and Weighted Content match Content-Based at launch** because no real popularity or interaction data exists yet. They are the natural next step.
 6. **Each additional signal meaningfully improves at least one metric:**
 - Workouts: Random → Popular: +0.04 precision, -0.380 diversity (bad trade)
 - Workouts: Popular → Content-Based: +0.02 precision, +0.480 diversity (clear win)
 - Meals: Content-Based beats Most Popular by **+0.48 precision** (0.780 vs 0.300)
-

Seminar Discussion Questions

Workout-specific

1. **Why does CF score 0?** → No interaction data exists; cosine similarity is undefined on empty vectors → cold-start failure demonstrated empirically
2. **Why is precision 1.0 not suspicious?** → Relevance is defined by the same goal-type weights the content scorer uses; alignment is intentional and correct
3. **Why is diversity important here?** → A fitness app serving 10 users the same list defeats the purpose; diversity=0.13 for Most Popular is a design failure

Meal-specific

4. **Why is precision the same (0.714) across all meal methods?** → Spoonacular pre-filters by diet type; the API does the hardest job. Look at coverage for the real performance difference.
5. **Why is CF dangerous for meals?** → It cannot enforce allergy/intolerance hard constraints — recommending a nut-based meal to someone with a nut allergy is a safety, not just a relevance, failure.

Cross-domain

6. **What happens to Hybrid/Weighted when Revfit grows?** → As interaction history accumulates, the popularity component gains signal — Hybrid's α lets us gradually shift from content-only to balanced personalisation
7. **What is the ground truth problem?** → Relevance here is preference-based (goal type match), not completion-based. Real ground truth requires A/B testing with actual user completion and feedback data.